

Parallel Applications

Laboratory 4

Deadline: 2 weeks

This assignment will help us learn and evaluate

- (i) two ways of writing parallel applications – using threads and using processes,
- (ii) two ways of doing inter-process communication – using shared memory and using pipes,
- (iii) two ways of doing process synchronization – using atomic operations and using semaphores.

Part I: Baseline application without any parallelism

This part is easy. Use your submission from laboratory 1. Update the makefile as indicated in the instructions below.

Part II

Now suppose you have a processor with three cores. You want to use all of them to make your application finish faster. You can do this by having the first core do *S1*. As pixels get ready, they are passed to the second core (don't wait for the *S1* of the whole image to be completed before communicating to *S2*). The second core does *S2* and passes its results to the third core. The third core does *S3* and the file writing. Do this in the following ways:

1. *S1*, *S2*, and *S3* are performed by 3 different processes that communicate via [pipes](#) (or [fifos](#)) ([further reading on pipes](#))
 2. *S1*, *S2*, and *S3* are performed by 3 different processes that communicate via [shared memory](#). Synchronization is done using [semaphores](#).
 3. *S1*, *S2*, and *S3* are performed by 3 different [threads](#) of the same process. They communicate through the process' address space itself. Synchronization is done using [atomic operations](#).
- Devise a method to prove in each parallel case that the pixels were received as sent, in the sent order. Describe the method in your report.
 - Study the run-time and speed-up of each of the approaches and discuss.
 - It is likely that the file reading and writing times dominate, and so the speed-up obtained by using three cores is negligible. So we will modify our experiment to make it compute-intensive, instead of IO-intensive. Read the image only once, but perform the transformation 1000 times. That is, *S1*, *S2*, and *S3* are done 1000 times each. The results of the first run of *S1* are used by the first run of *S2*, and the results of the first run of *S2* are used by the first run of *S3*. For the

second run of S1, it uses the same input image that has already been read into memory. The results of the second run of S1 are used by the second run of S2, and so on. The results of the 1000th run of S3 are written to the file.

- Discuss the relative ease/ difficulty of implementing/ debugging each approach.

Submit a single zip file with the source code (organized into multiple folders, one for each question), a makefile, an input *ppm* image, and a report.

- *make part1* should compile the Part I version of the code and run it, creating the file *output_part1.ppm*
- *make part2_1* should compile the multi-process, pipe version of the code and run it, creating the file *output_part2_1.ppm*
- *make part2_2* should compile the multi-process, atomic operation version of the code and run it, creating the file *output_part2_2.ppm*
- *make part2_3* should compile the multi-thread, semaphore version of the code and run it, creating the file *output_part2_3.ppm*